

Provenance-Certified KKR Convergence Assistance: ml_assist, PotentialBank v1, and Same-Structure Warm Starts

An AiiDA-native early-abort framework and a certified potential-reuse bank —
with a live pilot, a held-out benchmark, and explicitly benchmarked limits

Mohammad Hemmati*

¹Peter Grünberg Institut and Institute for Advanced Simulation, Forschungszentrum Jülich and JARA, D-52425 Jülich, G

²RWTH Aachen University, D-52062 Aachen, Germany

2026-08-13

Abstract

Self-consistent-field (SCF) convergence is the practical bottleneck of KKR and Bogoliubov-de Gennes (BdG) Green’s-function DFT campaigns: many runs stall or diverge, wasting cluster time. We present an **AiiDA-native convergence-control framework**, `ml_assist`, that scores a run’s first ten SCF iterations and can abort doomed runs early — an opt-in layer with a zero-risk default. The framework ships a **frozen numpy random-forest scorer** (the AiiDA daemon has no scikit-learn) alongside a **simple RMS-at-iteration-10 baseline**, both wrapped in a **per-cohort conformal early-abort policy** that controls the false-abort rate after per-cohort calibration under calibration/test exchangeability, with Mondrian per-stratum thresholds for conditional coverage. A **live Fe-doped NbSe₂ pilot validates the abort machinery end-to-end**: a true early abort saved ~ 130 SCF iterations against its diverging counterfactual, and a would-be false abort on a slowly-converging run was correctly withheld ($N=5$, descriptive). On a **336-run, leakage-clean held-out replay with per-run convergence-target (QBOUND) labels**, the learned model does *not* beat the RMS@10 baseline on the primary NbSe₂ normal-state cohort (held-out AUC 0.83 vs. 0.90), *does not transfer across materials* (other-material AUC $\approx$ chance), and a NbSe₂-seeded conformal threshold does not transport (realized false-abort 0.47) — though **per-cohort recalibration restores** realized $\approx$ target α . We conclude that *reliable deployment requires per-cohort calibration and strong simple baselines*; the honest, provenance-tracked framework — not model superiority — is the contribution. We then present **PotentialBank v1**, a provenance-certified extension of AiiDA-KKR that realizes the warm-start lever: it retrieves, certifies, and reuses compatible *same-structure* converged potentials. AiiDA-KKR already provides the restart and `potential_overwrite` primitives; our contribution is the **certification and retrieval layer above them** — byte-exact potential integrity (three-way SHA-256), exact Z-sequence/NATYP compatibility, non-CPA/non-BdG scope control, and no guessed mapping. PotentialBank v1 is frozen at **237 unique certified donors (238 nodes)** across **8 material families** and **33 F2-eligible materials**, with same-structure warm starts cutting SCF iterations substantially (per-family, mixing- and contour-qualified). We claim *no cross-material transfer and no guaranteed acceleration*: donors reuse the identical structure, and nine hard cases are documented with root causes. Separately, an engineering effort repaired four cell-inappropriate BdG-workchain defaults, enabling the first

*Corresponding author: mohammad.hemmati@rwth-aachen.de. ORCID: 0009-0006-4205-5706.

end-to-end multi-atom BdG run and its memory-scaling number. Everything ships as an installable `aiida-kkr-mlassist` extension (with an optional convergence advisor), unit tests, and pre-registered success gates.

Keywords. KKR Green-function DFT; SCF convergence; AiiDA provenance; potential reuse / warm start; conformal prediction; reproducible workflows; certified donor potentials.

Program summary

Program title: `aiida-kkr-mlassist` / KKR Convergence Advisor / PotentialBank v1.

Developer’s repository: https://github.com/MRHemmati/kkr_ml_tools (CPC library link and DOI assigned on acceptance).

Licensing provisions: MIT (open source).

Programming language: Python (pure-`numpy` runtime core; AiiDA / `aiida-kkr` used only for submission).

Nature of problem: KKR SCF convergence wastes HPC time when runs stall or diverge; restarting from a converged potential is standard and AiiDA-KKR provides the primitives, but choosing a *compatible, provenance-safe* donor is error-prone and unaudited.

Solution method: an AiiDA-native, opt-in `ml_assist` early-abort wrapper (pure-`numpy` scorer + RMS@10 baseline under a per-cohort conformal policy) and **PotentialBank v1**, a provenance-graph donor catalog admitting potentials only under strict E1 certification — byte-exact SHA-256 integrity, exact Z-sequence / NATYP compatibility, non-CPA / non-BdG scope, and no guessed mapping — with an optional same-structure warm-start advisor.

Restrictions: normal-state, non-CPA for PotentialBank v1; CPA / BdG and cross-material prediction are not supported; warm-start reuse is same-structure only, with no guaranteed acceleration; raw potentials remain in the AiiDA repository (referenced by SHA-256 / provenance).

Unusual features: strict provenance certification with byte-exact hashing, a no-guessed-mapping rule, per-family same-structure warm-start benchmarks, and a documented catalog of hard cases.

Additional comments: PotentialBank v1 is frozen at 237 unique certified donors (238 nodes) across 8 families, 33 F2-eligible materials.

Contents

1	Goals and scope	3
2	Background and tooling (what we used and why)	4
3	Data and the provenance-mining audit	4
4	The convergence classifier	5
5	The held-out benchmark: does the model beat a simple RMS rule?	6
6	The conformal early-abort policy	8
7	Deployment evaluation (per-cohort, calibrated)	9
8	Live pilot: the abort machinery works end-to-end	10
9	The BdG Δ -channel	11
10	Workflow hardening: four BdG footguns and the memory number	11
11	Software package	12

12 Pre-registration and success gates (superseded)	13
13 Side studies and honest negatives	13
14 Reproducibility and governance	13
15 PotentialBank v1: L3 warm-start realized	14
15.1 Related work	15
16 Status and roadmap	15
17 Software and data availability	16
18 Reproducibility	16
A Result tables	17
B Model provenance	17

1 Goals and scope

The parent project is a multi-year KKR / KKR-BdG superconductivity effort on 2H-NbSe₂ with transition-metal dopants (V/Mn/Fe/Co), plus a Bi/NbBi subset. SCF convergence is the practical bottleneck: many runs stall, diverge, or converge slowly, wasting cluster time. `ml_assist` targets three *acceleration levers*:

- L1 — early abort** stop runs that are predicted not to converge, recovering their wasted tail.
- L2 — probe-and-commit mixing** race a few mixing settings briefly, commit the best.
- L3 — warm start** reuse a converged potential to halve iterations.

This report is in three parts. **Part 1** (Secs. 1–8) covers `ml_assist` — the early-abort framework (L1), its conformal policy, the live pilot, and the leakage-clean held-out replay. **Part 2** (Sec. 15) presents **PotentialBank v1**, the realized warm-start lever (L3): a certified catalog of single-structure converged potentials with an optional convergence advisor and measured same-structure savings. **Part 3** (Sec. 16 onward) gives limitations and roadmap. The probe-and-commit mixing search (L2) reuses the same infrastructure and remains estimated by the pre-registered campaigns.

Positioning. KKR restart and potential reuse are *not new*: AiiDA-KKR [1] already exposes the low-level primitives — a KKR calculation can continue from a parent remote folder, automatically extracting the converged potential, and the Voronoi step accepts an explicit `potential_overwrite`. Our contribution is the **layer above** these primitives. We introduce a provenance-certified extension of AiiDA-KKR that systematically retrieves, validates, and reuses compatible **same-structure** donor potentials for reproducible KKR warm starts, supplemented by an optional convergence advisor whose limits are explicitly benchmarked. Concretely: a provenance-graph procedure discovers candidate donors; each is admitted only under a strict eligibility policy (byte-exact potential integrity via a three-way SHA-256 match, exact *Z*-sequence / NATYP compatibility, non-CPA and non-BdG scope control, and no guessed mapping); and the resulting same-structure warm starts are benchmarked, with hard cases and negative results reported honestly. This turns an available mechanism into an auditable, policy-controlled reuse system.

Design principles. (1) *Opt-in, zero-risk default*: absent an `ml_assist` input the workflow behaves identically to stock. (2) *Advisory by default*; abort only when explicitly enabled and the statistical guarantee is certifiable. (3) *Per-system deployment*: we do not claim cross-material transfer. (4) *Honesty first*: negative results are reported as-is and are publishable.

2 Background and tooling (what we used and why)

JuKKR (KKRhost, voronoi, KKRimp). The Green’s-function DFT engine. SCF proceeds by iterating the potential; each iteration reports an RMS error and charge-neutrality. *Why it matters for ML*: these per-iteration quantities are the raw signal from which convergence fate is predicted.

BdG extension. Superconducting pairing is treated by doubling the Hamiltonian into Nambu (particle–hole) space; KKRhost-BdG additionally prints a pairing amplitude Δ per iteration. *Why*: the pairing channel is a BdG-specific failure mode invisible to normal-state features.

AiiDA 2.7. Workflow engine and provenance store. All calculations, inputs, and outputs are nodes in a directed provenance graph. *Why*: it gave us a mineable corpus of historical runs *and* is the natural place to deploy an opt-in wrapper workflow with full audit trail.

aiida-kk. The AiiDA plugin providing `kk_scf_wc` (normal SCF) and `kk_bdg_wc` (BdG chain). *Why*: `ml_assist` subclasses `kk_scf_wc` to add the prediction hook.

scikit-learn (training only). Random forests for the classifier. *Why RF*: in distribution it beats logistic regression and the physical extrapolation heuristic and needs no feature scaling (Table 3). *But out-of-sample* the strong baseline to beat is a plain RMS-at-iteration-10 rule, which the model does *not* beat on the primary cohort (Sec. 5). *Constraint*: scikit-learn is *not* installed in the AiiDA daemon environment.

numpy (inference). We export the trained forest to a flattened-array `.npz` and traverse it in pure `numpy`, reproducing scikit-learn’s `predict_proba` exactly. *Why*: the daemon must score runs without scikit-learn; a `numpy`-only path removes that dependency (Sec. 4).

Conformal prediction. A distribution-free wrapper that turns a score into a decision with a controlled error rate. *Why*: it lets us *control the false-abort rate* $\leq \alpha$ after per-cohort calibration under calibration/test exchangeability, regardless of the model’s calibration — the safety property an abort policy needs (Sec. 6).

3 Data and the provenance-mining audit

The AiiDA database holds $\sim 7,700$ calculation nodes and $\sim 6,400$ workflows (2021–2026). We mined per-iteration RMS/charge-neutrality trajectories and outcomes. Three *leakage-class* corrections were essential and define the trustworthiness of every downstream number:

1. **Segment de-inflation.** Training rows were KKR *segments* (up to 5 per workflow), not runs. Deduplicating to one cold-start row per run reduced 3,405 segments to **2,710 runs** (1,851 normal, 90% converged).¹

¹Modeling uses the 1,830 normal runs with at least $T_{\text{obs}}+2$ iterations (21 too-short runs excluded, needed to form the observation window); the BdG Δ study uses the 641 of 846 BdG runs whose `out_kkr` was retrievable and parseable (Sec. 9).

2. **Censoring scope.** 389 excepted/killed calculations are 100% censored; the dataset is *completed runs only*. Crashes/NaN-kills are out of scope and handled by a watchdog, not the model.
3. **The Replay-OOF Invariant.** Any replay, preview, or “deployment simulation” number is out-of-sample or prospective *by construction*, never in-sample. This was adopted after an in-sample shadow replay produced a spurious recall of 1.000; the honest out-of-sample number is 0.70 (Sec. 7). The model that scores a run never trained on it (verified by node id).

These corrections deflated some earlier claims (e.g. a cross-material transfer number collapsed from 0.94 to chance) — an outcome we treat as a contribution, not an embarrassment.

4 The convergence classifier




Figure 1: Why early prediction is possible. Converging, stalling, and diverging runs already trace distinct shapes within the first ten iterations (shaded) — the only data the model sees.

Features. A 15-dimensional vector computed from the first $T_{\text{obs}} = 10$ iterations: log-RMS window statistics (first, last, mean, std, min, drop, slope), two charge-neutrality statistics, and three mixing parameters each with a missing-flag. A single canonical implementation is shared by training and inference so they cannot drift (verified bit-identical).

Model and in-distribution accuracy. A 400-tree balanced random forest. Its *in-distribution* 5-fold cross-validation AUC is strong where data is plentiful and honestly weaker where it is thin (Table 2); the pooled normal-state 5-fold AUC is 0.944. In distribution it beats logistic regression and a physics extrapolation heuristic (Fig. 2, Table 3). **We deliberately do *not* treat 0.944 as the deployment headline:** it is a same-distribution cross-validation number. The deployment-relevant question — how the frozen model scores runs it has never seen, against a simple RMS baseline — is answered by the leakage-clean held-out benchmark of Sec. 5, and there the model does *not* beat RMS@10 and does *not* transfer across materials. The model is therefore deployed per-system with per-cohort calibration.




Figure 2: *In-distribution only*. From ten iterations, the random forest cross-validates better than a physics extrapolation heuristic or logistic regression (AUC; 0.5 is a coin flip). This is *not* the deployment metric — see the held-out benchmark (Fig. 3), where a simple RMS@10 rule matches or beats the model out-of-sample.

A dropped complication (reported for the record). A schedule-aligned observation window was a pre-registered design lock intended to protect “flat-start” runs (flat, high-RMS early, then converging). On the deduplicated data it gave *no* measurable benefit: the fixed window already separates flat-starters at AUC ~ 0.90 (the discriminative signal is in charge-neutrality and mixing, not the flat RMS slope). It was dropped for v1; the frozen model uses the plain fixed window.

Pure-numpy export and the float32 gotcha. The forest is flattened to node arrays (children, feature, threshold, per-leaf class probabilities) and traversed in `numpy`. Golden-parity against scikit-learn is **exactly** 0.0, verified running under the daemon interpreter (`numpy` 1.26, no scikit-learn). One subtlety was decisive: scikit-learn casts inputs to `float32` for tree traversal, so a `float64` comparison mis-routes samples sitting on a threshold tie (an initial 10^{-2} parity failure). The `numpy` path compares in `float32`; embedded golden vectors (including the exact edge sample that exposed the bug) guard against regression.

5 The held-out benchmark: does the model beat a simple RMS rule?

The honest test of deployment value is a *leakage-clean, out-of-sample* evaluation with correct labels. We assembled **336 completed ifflurm runs verified absent from training** (by node id) and scored each run’s frozen first-ten-iteration features. Three corrections over an earlier pooled analysis were decisive (they revise our own earlier, over-optimistic numbers):

1. **Per-run QBOUND labels.** Convergence is judged against each run’s own convergence target, not a fixed 10^{-3} ; a fixed cut mislabelled 57 runs (many NbSe₂ runs use QBOUND=0.008).
2. **Censoring.** 15 excepted/killed runs (exit $\neq 0$) are dropped from all denominators.
3. **No pooled headline.** Numbers are reported *per cohort*; a single pooled AUC conflates a warm/restart population (which mostly converges) with cold starts (which mostly fail) and is not a deployment metric.

The result (Table 1, Fig. 3). On the primary target — NbSe₂-family normal-state — the frozen model reaches held-out AUC **0.83** [0.77, 0.89] (cold starts 0.79), while a **conformally-calibrated RMS@10 rule reaches 0.90** [0.85, 0.94] and gives *higher* failed-run recall and compute savings at matched false-abort at every α (Sec. 7). **The learned model does not beat the simple RMS baseline on the primary cohort.** It is anti-predictive across materials (other-KKR AUC 0.48; a warm/restart cohort 0.35) — a genuine transfer failure, since RMS@10 remains predictive on the same runs (0.90, 0.61). The one cohort where the model edges RMS@10 (“unknown”, 0.80 vs. 0.78) is small ($n=42$). Excluding the 56 borderline runs (final RMS within a factor 3 of QBOUND) leaves the ranking unchanged.

Table 1: Held-out benchmark, per cohort (leakage-clean, per-run QBOUND labels, censored dropped). AUC [95% bootstrap CI]. *No pooled headline number is reported.* A conformally-calibrated RMS@10 rule matches or beats the frozen model everywhere except the small “unknown” cohort; the model is anti-predictive across materials (other-KKR, warm/restart).

cohort	n	conv.	frozen ML	RMS@10	RMS-slope
NbSe₂-normal (primary, all)	167	38	0.831 [.77,.89]	0.897 [.85,.94]	0.740
— cold starts	128	12	0.785 [.68,.88]	0.901 [.80,.97]	0.568
— warm/restart	39	26	0.430	0.414	0.467
Bi/NbBi	28	21	0.898	0.939	0.816
other-KKR (cross-material)	45	34	0.476	0.904	0.583
BdG-misfit (normal model on BdG)	78	41	0.581	0.813	0.631
unknown	42	10	0.800	0.778	0.878
warm/restart-like (all families)	93	67	0.350	0.606	0.738




Figure 3: Held-out per-cohort AUC (per-run QBOUND labels), frozen ML vs. RMS@10, with 95% bootstrap CIs. On the deployment target (NbSe₂-normal cold, leftmost) RMS@10 leads; the model is at or below chance on other-material and warm/restart cohorts. There is no single pooled headline number.

Scope consequences. (i) BdG runs must be scored with the BdG model, not the normal one (normal-model-on-BdG AUC 0.58). (ii) Warm/restart runs are *outside the current abort policy*:

their failures are rare and not separable from the early window (both methods $\approx$ chance), so they are excluded from abort and left to run. (iii) No cross-material transfer is claimed for the learned model.

6 The conformal early-abort policy

Guarantee. Given a target α , we abort a run iff its predicted convergence probability is below a threshold calibrated on the campaign’s converged runs. One-sided split conformal makes

$$\Pr(\text{abort} \mid \text{the run would actually converge}) \leq \alpha.$$

Certiability needs ~ 19 calibration points at $\alpha = 0.05$ (~ 9 at 0.10); a campaign seeds these from material-matched history and refreshes online. The guarantee holds under *exchangeability* of calibration and test runs; material-matched seeding plus online refresh is how we keep that assumption approximately valid as a campaign drifts; a pre-registered violation trip-wire reverts to advisory mode if the realized false-abort rate ever exceeds the budget.

The marginal-vs-conditional gap (found and fixed). The split-conformal guarantee is *marginal*. Verified in out-of-sample replay: a single (marginal) threshold gives overall false-abort 0.049 at $\alpha = 0.05$ but **0.086 conditional on the flat-starter stratum** — the very runs the dropped aligned window was meant to protect. The fix is **Mondrian (group-conditional) thresholds**: one threshold per stratum (flat-starter vs. other, derived directly from the feature vector, agreeing 100% with the trajectory definition). This restores conditional control to **0.047** (Fig. 4, Table 4). Both vehicles deploy Mondrian.




Figure 4: The safety fix. A single global threshold breaks the false-abort promise on the tricky “flat-start” runs (red); one threshold per subgroup (green) restores the target rate within the evaluated strata/cohort.

Seeded thresholds do not transport — recalibrate per cohort. The guarantee holds only under *exchangeability* of calibration and test runs, and the held-out data show this is not automatic. A threshold seeded on historical NbSe₂ converged scores, applied to the broader held-out cohort, gave a realized false-abort rate of **0.47** at $\alpha = 0.10$ — far above target: the seed does *not* transport across cohorts. The fix is **per-cohort recalibration**: re-deriving the threshold on the deployment

cohort’s own converged runs restores realized false-abort to $\approx$ target ($0.10 \rightarrow 0.098$ for the model, 0.100 for RMS@10; Fig. 5, nested-bootstrap split-conformal, model weights frozen). The conformal guarantee is thus a *per-cohort calibration* property, not a transported one; deployment must recalibrate on the target cohort, and the trip-wire reverts to advisory if the realized rate exceeds budget.




Figure 5: The conformal threshold is a per-cohort object. A NbSe₂-seeded threshold applied to held-out data breaks the false-abort promise (0.47 at $\alpha = 0.10$); per-cohort recalibration (weights unchanged) restores realized $\approx$ target for both the model and RMS@10.

Fallback discipline. For BdG, if the Δ channel is unparseable the policy falls back to an rms-only model *with its own* conformal threshold (never comparing an rms-only score to an rms+ Δ threshold); both thresholds are certified independently.

7 Deployment evaluation (per-cohort, calibrated)

All numbers are out-of-sample (Replay-OOF Invariant), accounted in SCF iterations (hardware-invariant), and evaluated *on the primary NbSe₂-normal cohort with per-cohort recalibration* (Sec. 5). Each abort is charged at the deployment horizon (cumulative iteration 20, when the model first scores under `nsteps=10`).

Recovery and false-abort, ML vs. RMS@10 (Fig. 6, Table 5). With per-cohort recalibration both policies hold realized false-abort at target ($\alpha = 0.05/0.10/0.20 \rightarrow 0.04/0.10/0.20$). At matched α the **RMS@10 baseline recovers more wasted compute and catches more doomed runs than the model**: failed-run recall 0.64/0.78/0.86 (RMS@10) vs. 0.53/0.62/0.71 (ML), and fraction of wasted iterations saved 0.61/0.73/0.81 vs. 0.54/0.62/0.69. The model’s abort precision stays high (≥ 0.96), but it is the weaker recovery policy on the primary cohort. This is the deployment consequence of the held-out benchmark: *a simple RMS@10 rule, per-cohort calibrated, is the stronger early-abort policy here.*

α economics. Because a false abort under the auto-retry ladder *warm-restarts* (it does not discard a converging run), its cost is small, which pushes the net-optimal α upward. The framework ships two profiles — $\alpha = 0.05$ (conservative) and $\alpha = 0.10$ (throughput) — but the choice of *scorer*




Figure 6: Calibrated risk-savings on the primary NbSe₂-normal cohort (per-cohort recalibration, weights frozen). (a) failed-run recall and (b) fraction of wasted iterations saved, both at matched realized false-abort. RMS@10 (blue) dominates the frozen model (red) at every α .

(RMS@10 vs. model) is now itself a per-cohort decision, with RMS@10 the default on the primary cohort.

The single-material prospective pilot (narrow). An earlier prospective check calibrated on 329 historical NbSe₂ converged scores and applied to a held-out 2026 NbSe₂ grid (20 runs, verified absent from training) gave a realized false-abort rate of 0.000 at $\alpha = 0.05$ and 0.10. **This is a narrow, single-material, in-distribution result** and must not be read as broad safety: the same seed, applied to the wider held-out cohort, gave false-abort 0.47 (Sec. 6, Fig. 5). It demonstrates that *material-matched* seeding worked on that single 2026 NbSe₂ grid, and reinforces the per-cohort-calibration requirement — not cross-cohort transport.

8 Live pilot: the abort machinery works end-to-end

Beyond retrospective replay, we ran the policy *live* in the AiiDA daemon through a staged, gated Fe-doped NbSe₂ pilot (Gate 0 → read-only pre-flight → advisory smokes → gated package install/-validation → stock baselines → ml-assist arm). Five systems, each paired against a matched stock counterfactual (same structure, cold start, mixing, segment cadence, partition). The pilot validates the *machinery*, not generalization; at $N=5$ every effect size is descriptive.

What the pilot shows. (i) A **true early abort** (S5) fired live end-to-end — hook → numpy inference → Mondrian per-stratum threshold → decision → telemetry → exit 701 — and saved ≈ 130 iterations. (ii) A **would-be false abort was correctly withheld** (S1, advisory), confirming the conservative-advisory default is the right guard for slowly-converging runs. (iii) An **honest miss** (S3): a run whose stall emerges late was not caught by the early window — consistent with the held-out finding that late stalls are the hard case. The abort mechanism, telemetry, and paired accounting are validated; the $N=5$ deltas are descriptive, not a powered savings claim.




Figure 7: Live pilot exemplars (schematic of recorded runs). **S5** (Fe 92%) diverges; the frozen model scored $P(\text{converge}) = 0.297$ below the abort threshold and issued a live `ERROR_MLASSIST_ABORT` (exit 701) at cumulative iteration 20, saving ≈ 130 SCF iterations vs. its 150-iteration diverging counterfactual. **S1** (pristine) scored low ($P = 0.262$) yet converges; run advisory, its would-be abort was correctly *withheld* — the deliberate conservative trade-off.

9 The BdG Δ -channel

The standard parser exposes RMS and charge-neutrality but not the pairing amplitude. We wrote a parser for the `1-independent delta` lines KKRhost-BdG prints (two spin components per iteration, before each iteration marker), yielding a per-iteration Δ trajectory (length verified equal to the iteration count). Adding seven Δ -shape features lifts within-BdG *in-distribution* AUC from **0.946 to 0.957** (95% CI on the lift $[+0.000, +0.023]$; Δ -only AUC 0.884). (These are in-distribution 5-fold numbers, like Table 2, not held-out; and BdG runs must be scored with this BdG model — the normal model on BdG data reaches only AUC 0.58, Sec. 5.) The lift is modest here because the historical corpus barely contains pairing-*instability* runs (where Δ grows) — precisely the regime the parser targets for Vehicle A — so it is expected to under-state the value on the campaigns it is built for. The BdG model carries Δ ; every Vehicle-A run logs both the rms-only and rms+ Δ scores to measure the lift prospectively.

10 Workflow hardening: four BdG footguns and the memory number

A separate but essential outcome: making `kk_r_bdg_wc` run a multi-atom cell end-to-end. (Much historical BdG ran instead as standalone `KkrCalculations` in groups; the workchain is the path for *new* BdG runs, and the Δ -parser of Sec. 9 works on any BdG `out_kkr` regardless of how it was launched.) Four defaults, all tuned for a single-atom bulk reference, broke on an 8-atom NbSe₂ cell; each was fixed and validated, and the fixes are folded into the workchain (v0.2.1–v0.2.4) and mirrored to the working git branch:

1. **Multi-grid ranks.** MPI ranks must divide *every* step’s energy-point grid; a chain with a 24-point normal step and a 32-point BdG step requires a common divisor (or per-step ranks).

2. **Cell-aware RCLUSTZ.** The screening-cluster radius now inherits the cell’s value instead of a bulk default that overran the compiled cluster limit (NACLSD).
3. **NSPIN-aware spin coupling.** BdG with NSPIN=2 requires coupled spin channels; the workchain now drops the incompatible decoupling flag automatically.
4. **Cell-aware BZDIVIDE.** The k -mesh inherits the cell’s value; a bulk-dense mesh overran the compiled k -point limit (KPOIBZ) once BdG lowered the symmetry.

With these fixed, an 8-atom BdG chain completed for the first time, giving the **Vehicle-A memory number**: BdG job memory is $\approx 1.0 \text{ GB/rank}$, *flat with rank count* (energy-parallel replication), i.e. $2.32\times$ the same-cell normal-state memory — a mesh-robust scaling factor used for cluster packing (Fig. 8).




Figure 8: The BdG cost. Doubling the Hamiltonian into Nambu space raises memory per core by $\approx 2.3\times$; the factor is stable across core counts, which sets how many jobs fit on a node.

11 Software package

Everything ships as an installable `aiida-kkr-mlassist` extension (src-layout), with the pure-numpy core cleanly separated from the AiiDA workchain so inference needs no AiiDA and the daemon needs no scikit-learn:

```
aiida_kkr_mlassist/
  features.py canonical 15-D feature schema (train == infer)
  infer.py MlassistModel - numpy random forest (sklearn-exact)
  conformal.py ConformalAbortPolicy + MondrianConformalAbortPolicy
  ranks.py multi-step MPI-rank / energy-grid validator
  bdg_parser.py BdG Delta-channel parser + features
  decision.py config, decision state-machine, telemetry, graceful failure
  workchain.py MlassistKkrScfWorkChain (entry point kkr.mlassist)
  models/ frozen artifacts + manifests + MODEL_CARD.md
```

The workchain is opt-in (an absent input reproduces stock behaviour) and records every decision in node extras (advisory shadow data and audit trail). A unit-test suite (feature parity, golden parity, rank validation, conformal shadow-replay, Mondrian control, decision state-machine, BdG parser, dual-scoring) runs under the daemon interpreter to prove the deployment path is scikit-learn-free.

12 Pre-registration and success gates (superseded)

Deployment was originally pre-registered as two campaigns (one shared preamble) with frozen model hashes, the estimand definitions (TOTAL_CAMPAIGN_ITER, MEDIAN_ITER_TO_CONV), an A/B fairness rule (interleaved within a partition, wall-clock never pooled across partitions), a numeric conformal-violation trip-wire, and a pre-commitment that failures are reported as-is. **This original plan — in particular the Vehicle-B throughput campaign and its *G3* target — was superseded** and is retained here only for provenance:

Vehicle A (BdG, conservative). rms+ Δ model, dual-scoring telemetry, $\alpha = 0.05$, Mondrian thresholds; a numeric near-criticality criterion controls how many cells are abort-enabled so the safety rule cannot silently starve the recovery gate; *G2* target: recover $\geq 50\%$ of the stock arm’s own waste, zero conformal violations, physics-equality pass. (*Not run — BdG live campaign not launched.*)

Vehicle B (normal, throughput). $\alpha = 0.10$; *paired within-system* A/B (Wilcoxon on iterations) for power; a paired warm-and-cold subset; an explicit racing ladder, commit rule, and warm-start donor-compatibility checks; *G3* target: median iterations-to-converge and total campaign iterations both down $\geq 25\%$.

What actually happened. The full Vehicle-B campaign was *not* run: the planned dopant breadth (Mn/V/Co) lived on the CLAIIX cluster, which became unavailable, collapsing the design to an Fe-only iffslurm pilot. That pilot ($N=5$, Sec. 8) validated the abort *machinery* but is underpowered for a Wilcoxon effect-size test, so **G3 is not an active claim**. The held-out benchmark (Sec. 5) then reshaped the deployment recommendation itself: per-cohort calibration with an RMS@10 comparator, rather than a fixed $\geq 25\%$ throughput gate. Any future campaign must re-pre-register against that recommendation.

Physics-equality tolerances (empirical). Set from the intrinsic numerical reproducibility — the per-run total-energy variation over the last converged iterations (plateau noise, median 1.6×10^{-7} Ry/atom over 2,833 runs). Between-run “rerun” scatter proved unusable (formula-level matching still conflates dopant configurations). The adopted total-energy tolerance is $3 \times$ the plateau-noise 90th percentile, $\approx 3 \times 10^{-5}$ Ry/atom — neither contaminated nor pathologically tight.

13 Side studies and honest negatives

Impurity (KKRimp) feasibility. Deferred: 2,390 impurity workchains exist, but their stored RMS trajectory has median length 1 — per-iteration RMS is not in provenance — so early-window prediction cannot be priced without richer logging. A logging change, not a modeling gap.

Warm-start (L3) matched pairs. Underpowered: zero system-matched cold/warm converged pairs exist historically, so the causal warm-start effect is not retrospectively estimable and is deferred to the pre-registered randomized design.

14 Reproducibility and governance

The AiiDA database is treated read-only for analysis; any code-environment or database change is confirmation-gated with a printed plan. Frozen models are versioned artifacts with SHA-256 and

a manifest recording feature schema, training-data hash, library versions, and out-of-sample AUC (Appendix B). The Replay-OOF Invariant, the two-denominator reporting rule, and the negatives-are-publishable pre-commitment are the standing methodological safeguards.

15 PotentialBank v1: L3 warm-start realized

The L3 lever — reuse a converged potential instead of a cold Voronoi start — is realized as **PotentialBank v1**: a provenance-clean catalog of certified single-structure converged KKR potentials, each carrying a byte-exact compatibility fingerprint. A donor is admitted only if it passes *E1 certification* verbatim: the parser succeeds and byte-exactly round-trips, the potential’s per-type *Z*-sequence matches the structure site order (empty spheres as $Z=0$), NATYP equals the site count, the cell is non-CPA and non-BdG, NSPIN/SOC/LMAX/mesh parse, no mapping is guessed, and the staged, AiiDA-retrieved, and manifest SHA-256 all agree. CPA, disordered, and BdG potentials are *excluded by construction*.

Coverage. v1 contains **238 certified donor nodes = 237 unique-SHA donors** (one recorded byte-duplicate: E1 Bi₂X₄ pk 24009≡24011), spanning **8 families, 35 materials, 42 compatibility classes**, of which **33 materials are F2-eligible** (≥ 2 independent certified donors): NbSe₂/Bi (E1, 160 donors), simple metals Nb/Al/Cu/Ag/Mo (24), heavy-SOC Au/W/Pt/Sb/Pb (24), magnetic Fe/Ni/Co (13), ordered intermetallics CuZn/Ni₃Al/FeAl (6), covalent empty-sphere Si/Ge (4), the metallic layered chalcogenide TiSe₂ (2), and ionic rocksalt MgO/NaCl (4). Each family required its own certified settings recipe (screening cluster RCLUSTZ, Chebyshev log panel R.LOG/NPAN, BZ mesh, SOC key routing, explicit empty-sphere mapping, and a mixing-path variation to obtain independent F2 pairs for deterministic insulator SCF).

Warm-start savings (same-structure). Injecting a certified donor as the start potential cuts iterations-to-converge by large, reproducible margins, *donor-choice-independent* (33/33 targets: identical iterations from either donor) and direction-symmetric where cold-comparable:

family	cold iters	warm iters	% saved
simple metals	66–102	4–7	91–94 %
heavy-SOC	122–136	2–3	~97 %
magnetic	131–136	5–6	~96 %
covalent Si/Ge	117–119	2–3	97–98 %
intermetallics	128–152	5–7	95–97 %
TiSe ₂ (layered)	1188–1354	≤ 100	~92 %
ionic MgO/NaCl	171–750	2–500	33–99 %

The magnitude is set by the target’s mixing accelerator (a Broyden target polishes a warm start in a few iterations; a pure-linear target still crawls, e.g. 500 vs. 750 for MgO/NaCl), so each row is a *matched same-mixing* warm-vs-cold comparison and we report per-family, never a single pooled headline. Contour caveats are recorded (TiSe₂ is coarse-converged only, as the fine contour destabilizes the zero-gap semimetal). Coverage, F2-eligibility, warm-vs-cold, savings distribution, the parked-hard-case map, and the provenance workflow are in Figs. 1–6 of the v1 synthesis package.

Limitations. PotentialBank v1 supports **same-structure reuse only**; it is *not* a cross-material transfer claim and *not* a guaranteed acceleration (donor≡target geometry, chemistry, and settings; savings depend on structure/mixing/contour). CPA/doped/disordered and BdG are out of scope.

Nine hard cases are parked with root causes (MoS₂ and NbS₂ NACLSD, WSe₂ SOC contour/EMIN instability, TaS₂ SOC-NaN, TiAl VERTEX3D, Bi₂Te₃/Bi₂Se₃ hard-convergence, CaF₂/LiF). The warm-start outputs are benchmarks and are *not* themselves certified as donors. For `ml_assist`, the early-abort mechanism is validated live but broad ML transfer is *not* proven — it does not beat a simple RMS@10 baseline out-of-sample and requires per-cohort recalibration.

15.1 Related work

Workflow and provenance. We build directly on AiiDA-KKR / JuKKR [1], which provide the KKR restart and `potential_overwrite` primitives our warm start uses. `aiida-mlip` [2] is a sibling AiiDA plugin that applies the same provenance discipline to machine-learned interatomic potentials, a different target. **Predicting quantities to accelerate SCF.** A growing body of work *predicts* an electronic-structure quantity to speed up or bypass self-consistency: ChargeE3Net [7] predicts the charge density and reports a 26.7% reduction in SCF iterations when initializing DFT on *unseen* materials; SALTED [8] learns 3D electron densities via density fitting; and DeepH / DeepH-E3 [9, 10] predict the DFT Hamiltonian directly. `mp-pyrho` [6] provides representation-independent charge-density regridding and an associated density database. **Representations and MLIP tooling.** DDescribe [3] supplies atomistic descriptors; DP-GEN [4] performs concurrent-learning generation of deep-potential models; `dftio` [5] standardizes DFT outputs for machine learning. **Where we differ.** These methods *learn to predict* a density, potential, or Hamiltonian and, in the best cases, transfer across chemistries. PotentialBank v1 makes **no prediction and no cross-material claim**: it *reuses an existing, certified potential for the identical structure*. The novelty is the certification and retrieval policy — integrity, compatibility, scope, and auditability — not a new physical model or a learned mapping. We are not aware of a verified KKR-specific tool that provides this certified same-structure reuse layer.

What is and is not novel. We do **not** claim to have invented KKR restart, potential overwriting, or SCF warm starting — those primitives exist in AiiDA-KKR and are standard practice. We do **not** predict potentials, densities, or Hamiltonians, and we make no cross-material transfer claim. What is novel is the **provenance-certified reuse layer**: automated donor discovery over the provenance graph; a strict, auditable eligibility policy (byte-exact integrity, exact structural compatibility, non-CPA / non-BdG scope, no guessed mapping); an optional convergence advisor with explicitly benchmarked limits; and honest same-structure warm-start benchmarks together with a documented catalog of hard cases and negative results.

16 Status and roadmap

Complete and validated: the framework and its pure-numpy artifact; the conformal (Mondrian) abort policy; the live pilot (abort mechanism validated end-to-end, Sec. 8); the leakage-clean held-out benchmark (Sec. 5); the BdG Δ -parser and frozen BdG models; the `aiida-kkr-mlassist` package with tests; the BdG workchain hardening and memory number; the physics-equality tolerance; the pre-registrations. **Key finding for deployment:** the learned model does not beat a simple RMS@10 rule out-of-sample on the primary cohort, does not transfer across materials, and its conformal guarantee requires *per-cohort recalibration*. Deployment is therefore per-cohort and baseline-benchmarked: the default scorer on the primary cohort is RMS@10, and any campaign must recalibrate the abort threshold on its own converged runs and compare model vs. RMS@10 before enabling abort. **Remaining modeling gap:** late-stall detection (largely missed by the early

window, for both the model and RMS@10). **L3 (warm start) is delivered** as PotentialBank v1 (Sec. 15): 237 certified single-structure donors across 8 families with measured same-structure savings; **L2** (mixing search) remains estimated by the pre-registered campaigns. PotentialBank v1 is same-structure reuse only — no cross-material transfer, no guaranteed acceleration.

17 Software and data availability

The software is open source under the **MIT** license and is available at https://github.com/MRHemmati/kkr_ml_tools (a tagged release and archival DOI will be minted on acceptance and recorded here). The repository ships the installable `aiida-kkr-mlassist` package and the `kkr_convergence_advisor`, unit tests, the frozen model artifacts with SHA-256 manifests, the PotentialBank v1 manifest / tables / figures, the advisor index, and two walkthrough notebooks. **Data availability:** the derived artifacts (manifests, benchmark tables, model manifests) are included; the raw KKR potential files are *not* redistributed — they remain in the AiiDA provenance store and are referenced by byte-exact SHA-256 and provenance handles (`seg_pk`, `wc_pk`, `uuid`), so their identity is verifiable by anyone with access to that store. Live KKR submissions require the JuKKR codes and the authors’ AiiDA database and cluster environment. *Versioning:* the Python *package* (`aiida-kkr-mlassist`) and the *data release* (PotentialBank and advisor index) are versioned independently — the package is at 0.1.1, while the certified potential catalog and advisor index are the frozen *PotentialBank v1* release; these are distinct artifacts and the different version labels are intentional, not an inconsistency.

18 Reproducibility

Offline, without HPC or AiiDA: install from source (`pip install -e ., numpy-only core`), load the advisor index (`kkr_convergence_advisor.load_index()`), run the self-asserting unit tests (`PYTHONPATH=src python tests/*.py`), and execute the two walkthrough notebooks (the submission tutorial is safe-by-default and submits nothing). A lightweight smoke test reproduces the headline counts (237/238/8/33); `sha256sum -c reports/potentialbank_v1/MANIFEST.sha256` verifies the synthesis bundle. The reviewer-run checklist gives exact commands and expected outputs.

CRedit author contribution statement

Mohammad Hemmati: Conceptualization, Methodology, Software, Validation, Formal analysis, Investigation, Data curation, Writing — original draft, Writing — review & editing, Visualization.

Declaration of competing interest

The author declares no known competing financial interests or personal relationships that could have appeared to influence the work reported in this paper.

Acknowledgments

We thank the Bavarian Ministry of Economic Affairs, Regional Development and Energy for financial support within the High-Tech Agenda Project “Bausteine für das Quantencomputing auf Basis topologischer Materialien mit experimentellen und theoretischen Ansätzen”. Furthermore, this work was funded by the Deutsche Forschungsgemeinschaft (DFG, German Research Foundation)

under Germany’s Excellence Strategy — Cluster of Excellence Matter and Light for Quantum Computing (ML4Q) EXC 2004/1 — 390534769. *TODO: computing-time acknowledgment (this work’s KKR runs used iffslurm at Forschungszentrum Jülich — insert the correct allocation once known; do not reuse the JARA/CLAIX/JURECA wording from the earlier paper unless applicable).*

Use of AI assistance

AI coding and writing assistants were used to help draft text, generate code-edit suggestions, and organize reports; all scientific decisions, calculations, validation, and final claims were reviewed by the author.

A Result tables

Table 2: Per-material convergence AUC, *in-distribution* 5-fold cross-validation (deduplicated runs). Strong where data is plentiful; honestly weaker/underpowered where thin. **These are same-distribution CV numbers, not the deployment metric** — for the leakage-clean held-out result see Table 1. No cross-material transfer is claimed.

family	n	OOF AUC	95% CI
NbSe ₂	373	0.918	[0.86, 0.97]
Bi/NbBi	38	0.849	[0.69, 0.97]
mixed/unknown ²	1398	0.950	[0.92, 0.97]
pooled (normal, frozen model)	1830	0.944 ³	—

Table 3: The ML margin over baselines, *in-distribution only*. Out-of-sample the relevant baseline is RMS@10 (not the weak extrapolation heuristic), and it matches or beats the model per cohort (Table 1).

method	in-distribution AUC	held-out (primary cohort) AUC
physical extrapolation heuristic	0.549	—
logistic regression	0.826	—
RMS@10 rule (simple baseline)	—	0.897
random forest (ml_assist)	0.944	0.831

B Model provenance

Frozen artifacts (pure-numpy, golden-parity = 0 vs. scikit-learn, versioned by SHA-256):

³In-distribution 5-fold CV; *not* the deployment headline. Out-of-sample on the primary NbSe₂-normal cohort the frozen model reaches AUC 0.83 and is beaten by RMS@10 0.90 (Table 1).

²“mixed/unknown” = runs whose chemical formula could not be recovered from provenance and so could not be assigned to the NbSe₂ or Bi/NbBi families; they are the largest, cleanest group and anchor the pooled model.

Table 4: Marginal vs. Mondrian conformal, false-abort rate (out-of-sample). The marginal threshold under-protects the flat-starter stratum; per-stratum (Mondrian) thresholds restore conditional control.

α	marginal (overall)	marginal (flat-starters)	Mondrian (flat-starters)
0.05	0.049	0.086	0.047
0.10	0.098	0.164	0.093

Table 5: Calibrated early-abort on the primary NbSe₂-normal cohort (held-out, per-cohort recalibration, weights frozen; nested-bootstrap split-conformal, iterations). Realized false-abort holds at target for both; **RMS@10 recovers more and catches more doomed runs than the model at every α .**

α	policy	realized false-abort	failed-run recall	wasted-iters saved
0.05	frozen ML	0.036	0.53	0.54
0.05	RMS@10	0.039	0.64	0.61
0.10	frozen ML	0.096	0.62	0.62
0.10	RMS@10	0.100	0.78	0.73
0.20	frozen ML	0.196	0.71	0.69
0.20	RMS@10	0.186	0.86	0.81

model	role	features	in-dist. AUC ⁴
<code>mlassist_v1_normal_fixed10</code>	normal-state	15	0.944
<code>mlassist_v1_bdg_rmsdelta</code>	BdG primary	22	0.957
<code>mlassist_v1_bdg_rms</code>	BdG missing- Δ fallback	15	0.946

Each artifact ships a JSON manifest (feature names, training-data SHA-256, library versions, creation timestamp, golden-parity value) and a model card stating intended use, per-system scope, and the Replay-OOF Invariant.

References

- [1] P. Rüßmann, F. Bertoldo, and S. Blügel, “The AiiDA-KKR plugin and its application to high-throughput impurity embedding into a topological insulator,” *npj Comput. Mater.* **7**, 13 (2021), DOI 10.1038/s41524-020-00482-5. Software (JuKKR): <https://github.com/JuDFTeam/aiida-kkr>; docs <https://aiida-kkr.readthedocs.io> (User guide: Calculations, Workflows).
- [2] STFC, *aiida-mlip: machine-learned interatomic potentials in AiiDA* (software documentation, BSD-3). <https://github.com/stfc/aiida-mlip>; docs <https://stfc.github.io/aiida-mlip/>.
- [3] L. Himanen *et al.*, “DScribe: Library of descriptors for machine learning in materials science,” *Comput. Phys. Commun.* **247**, 106949 (2020); arXiv:1904.08875.

⁴In-distribution 5-fold CV. Held-out, the normal model scores 0.83 on its primary cohort (Table 1); weights were unchanged through all held-out and pilot evaluations.

- [4] Y. Zhang *et al.*, “DP-GEN: A concurrent learning platform for the generation of reliable deep learning based potential energy models,” *Comput. Phys. Commun.* **253**, 107206 (2020); arXiv:1910.12690.
- [5] DeePModeling community, *dftio: transcribing DFT outputs into ML-ready formats* (software documentation, MIT). <https://github.com/deepmodeling/dftio>; docs <https://deepmodeling.github.io/dftio/>.
- [6] J. Shen *et al.*, “A representation-independent electronic charge density database for crystalline materials,” *Sci. Data* **9**, 661 (2022); arXiv:2107.03540. Software: <https://github.com/materialsproject/pyrho>.
- [7] T. Koker *et al.*, “Higher-order equivariant neural networks for charge density prediction in materials,” *npj Comput. Mater.* **10**, 138 (2024); arXiv:2312.05388.
- [8] A. M. Lewis, A. Grisafi, M. Ceriotti, and M. Rossi, “Learning electron densities in the condensed phase” (SALTED), *J. Chem. Theory Comput.* **17**, 7203 (2021), DOI 10.1021/acs.jctc.1c00576; arXiv:2106.05364. Docs <https://salted.readthedocs.io>.
- [9] H. Li *et al.*, “Deep-learning density functional theory Hamiltonian for efficient ab initio electronic-structure calculation,” *Nat. Comput. Sci.* **2**, 367 (2022).
- [10] X. Gong *et al.*, “General framework for E(3)-equivariant neural network representation of density functional theory Hamiltonian” (DeepH-E3), *Nat. Commun.* **14**, 2848 (2023).